

# CSE 574 Introduction to Machine Learning

## Assignment 3 Building Neural Networks and CNN

### Team Members:

Naman Tejaswi - [namantej@buffalo.edu](mailto:namantej@buffalo.edu) (50435697)  
Sachin Sarsambi - [ssarsamb@buffalo.edu](mailto:ssarsamb@buffalo.edu) (50419943)

# Academic Integrity Statement

I/We certify that the code and data in this assignment were generated independently, using only the tools and resources defined in the course and that I/we did not receive any external help, coaching or contributions during the production of this work.

## Part I – Building a Basic NN

1. Provide brief details about the nature of your dataset. What is it about? What type of data are we encountering? How many entries and variables does the dataset comprise? Provide the main statistics about the entries of the dataset.

- The dataset being analyzed contains Income details of people across different countries and working class. The following are the columns and their corresponding datatypes:

Column Names	Data Type
age	int64
workclass	object
fnlwgt	int64
education	object
education.num	int64
marital.status	object
occupation	object
relationship	object
race	object
sex	object
capital.gain	int64
capital.loss	int64
hours.per.week	int64
native.country	object
income	object

- Here, there are 15 columns and 32561 data samples. The columns *age*, *fnlwgt*, *education.num*, *capital.gain* and *capital.loss* are of the type *int* whereas the columns of *workclass*, *education*, *marital.status*, *occupation*, *relationship*, *race*, *sex*, *native.country* and *income* are of the type *object*. The columns of the dtype *object* are later converted to *category* type. Using the correlation matrix it can be observed that there is a direct relationship between *income* and *hours per week*.

- The salient features of the dataset:

index	age	work class	fnlwgt	education	education. num	marital .status	occu patio n	relati onshi p	rac e	sex	capital.gain	capital.loss	hours.per. week	native. country	inc om e
cou nt	32561.0	3256 1	32561.0	3256	32561.0	32561	3256 1	32561	32 56 1	32 56 1	32561.0	32561.0	32561.0	32561	325 61
uni que	NaN	9	NaN	16	NaN	7	15	6	5	2	NaN	NaN	NaN	42	2
top	NaN	Priva te	NaN	HS- grad	NaN	Marrie d-civ- spouse	Prof- speci alty	Husba nd	W hit e	Ma le	NaN	NaN	NaN	United- States	<=5 0K
fre q	NaN	2269 6	NaN	1050 1	NaN	14976	4140	13193	27 81 6	21 79 0	NaN	NaN	NaN	29170	247 20
me an	38.581646 75532080	NaN	189778.36 651208500	NaN	10.080679 3403151	NaN	NaN	NaN	NaN	NaN	1077.6488 437087300	87.303829 734959	40.437455 852093000	NaN	NaN
std	13.640432 553581100	NaN	105549.97 769702200	NaN	2.5727203 320673400	NaN	NaN	NaN	NaN	NaN	7385.2920 84836110	402.96021 86495170	12.347428 681730800	NaN	NaN
mi n	17.0	NaN	12285.0	NaN	1.0	NaN	NaN	NaN	NaN	NaN	0.0	0.0	1.0	NaN	NaN
25 %	28.0	NaN	117827.0	NaN	9.0	NaN	NaN	NaN	NaN	NaN	0.0	0.0	40.0	NaN	NaN
50 %	37.0	NaN	178356.0	NaN	10.0	NaN	NaN	NaN	NaN	NaN	0.0	0.0	40.0	NaN	NaN
75 %	48.0	NaN	237051.0	NaN	12.0	NaN	NaN	NaN	NaN	NaN	0.0	0.0	45.0	NaN	NaN
ma x	90.0	NaN	1484705.0	NaN	16.0	NaN	NaN	NaN	NaN	NaN	99999.0	4356.0	99.0	NaN	NaN

- The head data of the income dataset:

```
df.head
<bound method NDFrame.head of      age workclass fnlwgt      education education.num    marital.status \
0     90      ?  77053  HS-grad       9        Widowed
1     82  Private  132870  HS-grad       9        Widowed
2     66      ? 186061 Some-college      10        Widowed
3     54  Private  140359  7th-8th       4       Divorced
4     41  Private  264663 Some-college      10      Separated
...   ...
32556  22  Private  310152 Some-college      10 Never-married
32557  27  Private  257302 Assoc-acdm      12 Married-civ-spouse
32558  40  Private  154374  HS-grad       9 Married-civ-spouse
32559  58  Private  151910  HS-grad       9        Widowed
32560  22  Private  201490  HS-grad       9 Never-married

      occupation relationship race   sex capital.gain \
0           ? Not-in-family White Female          0
1  Exec-managerial Not-in-family White Female          0
2           ? Unmarried Black Female          0
3  Machine-op-inspct Unmarried White Female          0
4  Prof-specialty Own-child White Female          0
...   ...
32556  Protective-serv Not-in-family White Male          0
32557  Tech-support   Wife White Female          0
32558  Machine-op-inspct Husband White Male          0
32559  Adm-clerical Unmarried White Female          0
32560  Adm-clerical Own-child White Male          0

      capital.loss hours.per.week native.country income
0            4356          40 United-States <=50K
1            4356          18 United-States <=50K
2            4356          40 United-States <=50K
3            3900          40 United-States <=50K
4            3900          40 United-States <=50K
...   ...
32556          0          40 United-States <=50K
32557          0          38 United-States <=50K
32558          0          40 United-States >50K
32559          0          40 United-States <=50K
32560          0          20 United-States <=50K

[32561 rows x 15 columns]>
```

- The column names and their corresponding information:

```
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   age         32561 non-null  int64   
 1   workclass   32561 non-null  object  
 2   fnlwgt     32561 non-null  int64   
 3   education   32561 non-null  object  
 4   education.num 32561 non-null  int64   
 5   marital.status 32561 non-null  object  
 6   occupation   32561 non-null  object  
 7   relationship 32561 non-null  object  
 8   race         32561 non-null  object  
 9   sex          32561 non-null  object  
 10  capital.gain 32561 non-null  int64   
 11  capital.loss 32561 non-null  int64   
 12  hours.per.week 32561 non-null  int64   
 13  native.country 32561 non-null  object  
 14  income       32561 non-null  object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

- The tail records of the income dataset:

```
df.tail
<bound method NDFrame.tail of      age workclass fnlwgt      education education.num    marital.status \
0      90      ?  77053  HS-grad       9      Widowed
1      82  Private  132870  HS-grad       9      Widowed
2      66      ?  186061  Some-college    10      Widowed
3      54  Private  140359   7th-8th       4      Divorced
4      41  Private  264663  Some-college    10      Separated
...
...      ...      ...
32556     22  Private  310152  Some-college    10  Never-married
32557     27  Private  257302  Assoc-acdm     12  Married-civ-spouse
32558     40  Private  154374  HS-grad       9  Married-civ-spouse
32559     58  Private  151910  HS-grad       9      Widowed
32560     22  Private  201490  HS-grad       9  Never-married

      occupation relationship race   sex capital.gain \
0      ? Not-in-family White Female          0
1  Exec-managerial Not-in-family White Female          0
2      ? Unmarried Black Female          0
3  Machine-op-inspect Unmarried White Female          0
4  Prof-specialty Own-child White Female          0
...
...      ...
32556  Protective-serv Not-in-family White Male          0
32557  Tech-support Wife White Female          0
32558  Machine-op-inspect Husband White Male          0
32559  Adm-clerical Unmarried White Female          0
32560  Adm-clerical Own-child White Male          0

  capital.loss hours.per.week native.country income
0        4356           40 United-States <=50K
1        4356           18 United-States <=50K
2        4356           40 United-States <=50K
3        3900           40 United-States <=50K
4        3900           40 United-States <=50K
...
...      ...
32556      0           40 United-States <=50K
32557      0           38 United-States <=50K
32558      0           40 United-States >50K
32559      0           40 United-States <=50K
32560      0           20 United-States <=50K

[32561 rows x 15 columns]>
```

- The shape of the dataset is:

```
df.shape
(32561, 15)
```

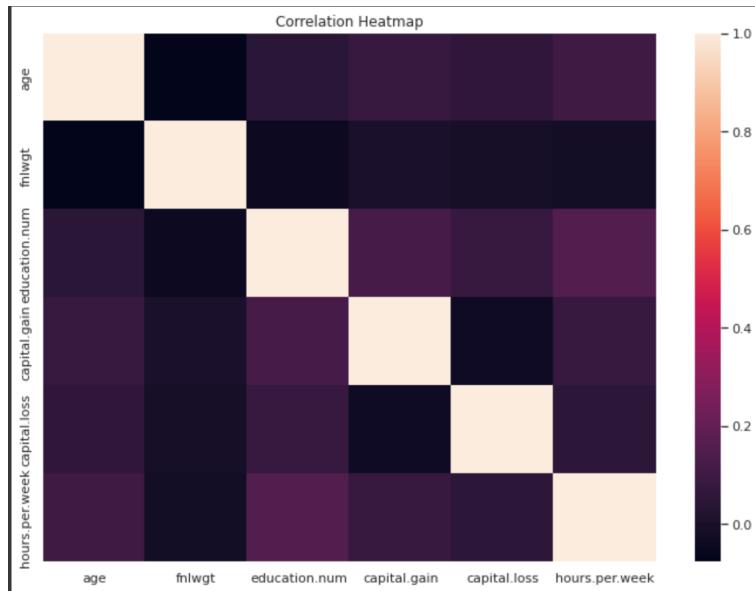
- The count of "?" character across various columns:

```
df[df=='?'].count()

age                  0
workclass            1836
fnlwgt                 0
education                0
education.num            0
marital.status            0
occupation            1843
relationship                0
race                  0
sex                  0
capital.gain                0
capital.loss                 0
hours.per.week                0
native.country            583
income                  0
dtype: int64
```

2. Provide at least 3 visualization graphs with short description for each graph.

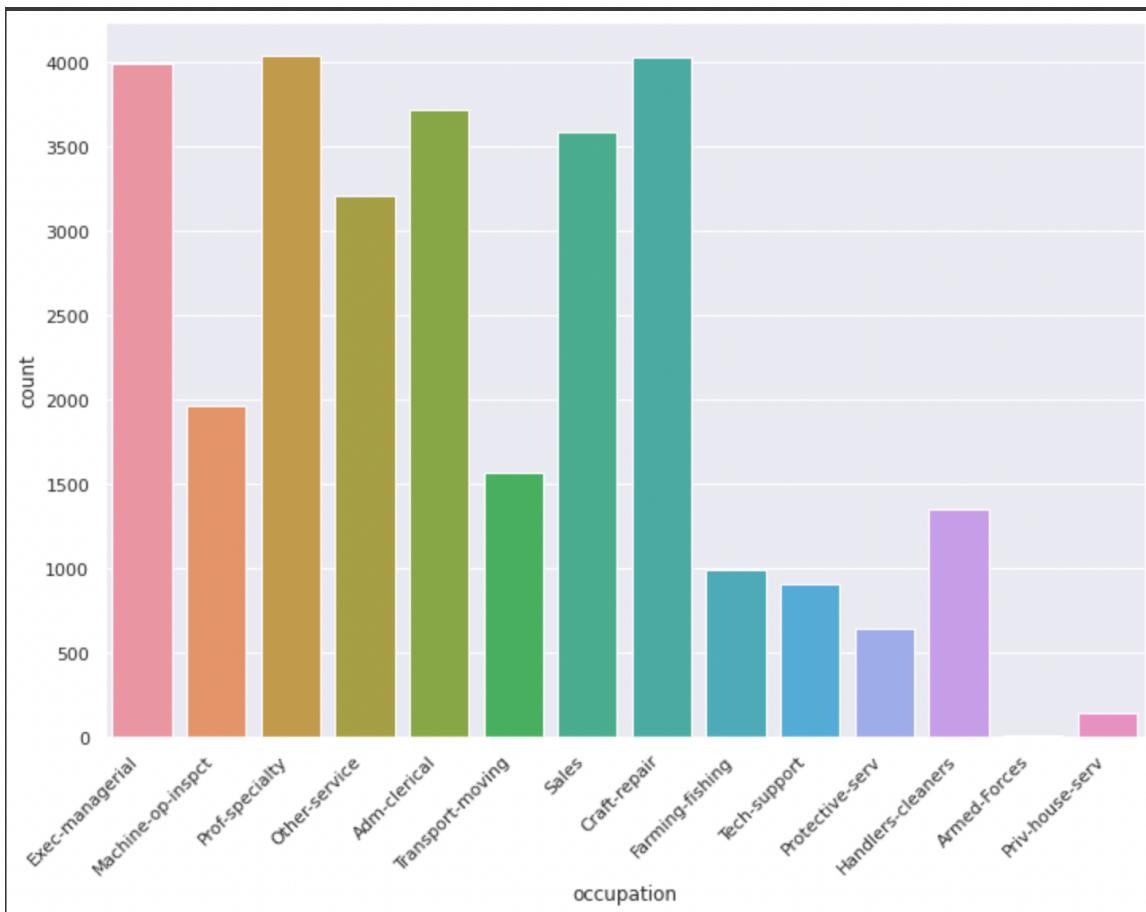
- *Correlation Matrix:*



df.corr()						
	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
age	1.000000	-0.076511	0.043526	0.080154	0.060165	0.101599
fnlwgt	-0.076511	1.000000	-0.044992	0.000422	-0.009750	-0.022886
education.num	0.043526	-0.044992	1.000000	0.124416	0.079646	0.152522
capital.gain	0.080154	0.000422	0.124416	1.000000	-0.032229	0.080432
capital.loss	0.060165	-0.009750	0.079646	-0.032229	1.000000	0.052417
hours.per.week	0.101599	-0.022886	0.152522	0.080432	0.052417	1.000000

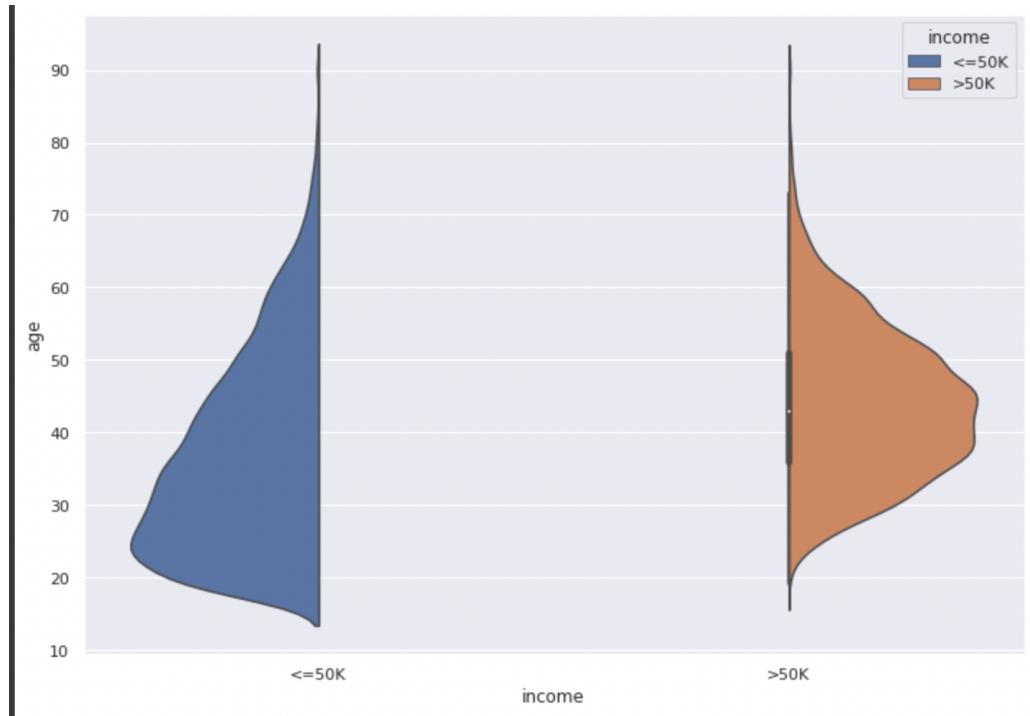
- We are using a heatmap to represent a correlation matrix between features *age*, *fnlwgt*, *education.num*, *capital.gain*, *capital.loss* and *hours.per.week*. A correlation matrix indicates relationship between various features.
- Here, the closest relation is between *hours per week* and *education column* with a correlation value of 0.15 and then followed by *education column* and *capital gain*.

- *Occupation Distribution:*



- The occupation with least number of employees is *Armed-Forces* (9 Employees) and the occupation with the greatest number of employees is *Prof-Specialty* (4140 Employees).

- *Violin Graph:*



- The above violin graph shows the distribution of income against employee age.
  - The density of employee with income  $\leq 50K$  is highest in the age group 20-30 years and the density of employees with income  $>50K$  is highest in the age group 35-45 years.
3. For the preprocessing part, discuss if you use any preprocessing tools, that helps to increase the accuracy of your model.
- *One-Hot Encoder:* To improve the accuracy, we include as many features as possible. To achieve this, we need to convert the category type data to numerical ones using the one-hot encode process. To apply one-hot encoding on the feature dataset, we use pandas in-built function `get_dummies` api.
  - *Min-Max-Normalization:* The `sklearn.preprocessing.MinMaxScaler` performs min-max normalization on the dataset so that the input features have the similar range of magnitude.

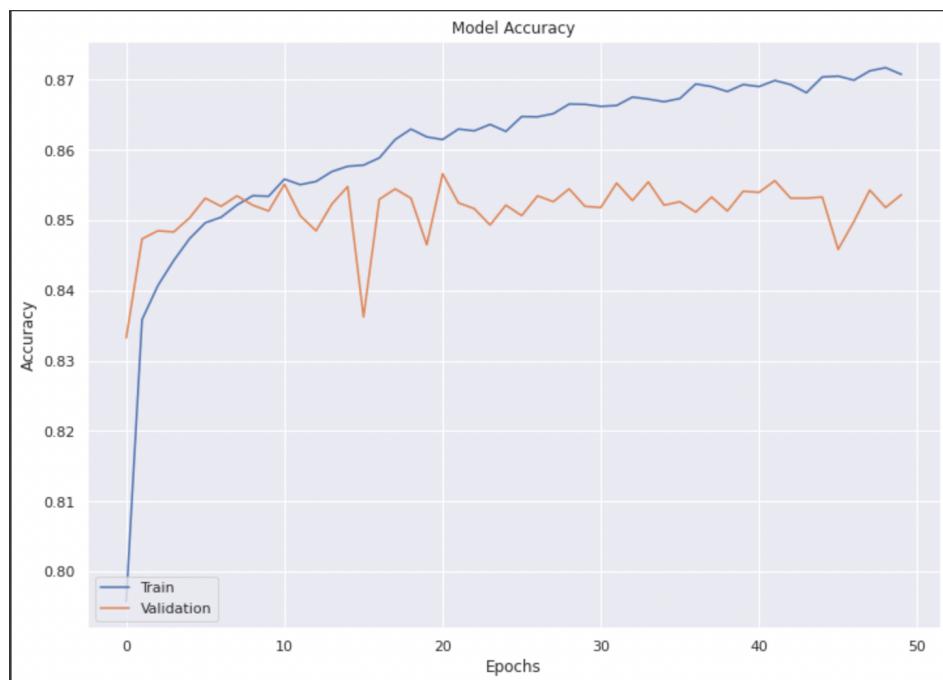
4. Provide the architecture structure of your NN

- The following is the model summary:

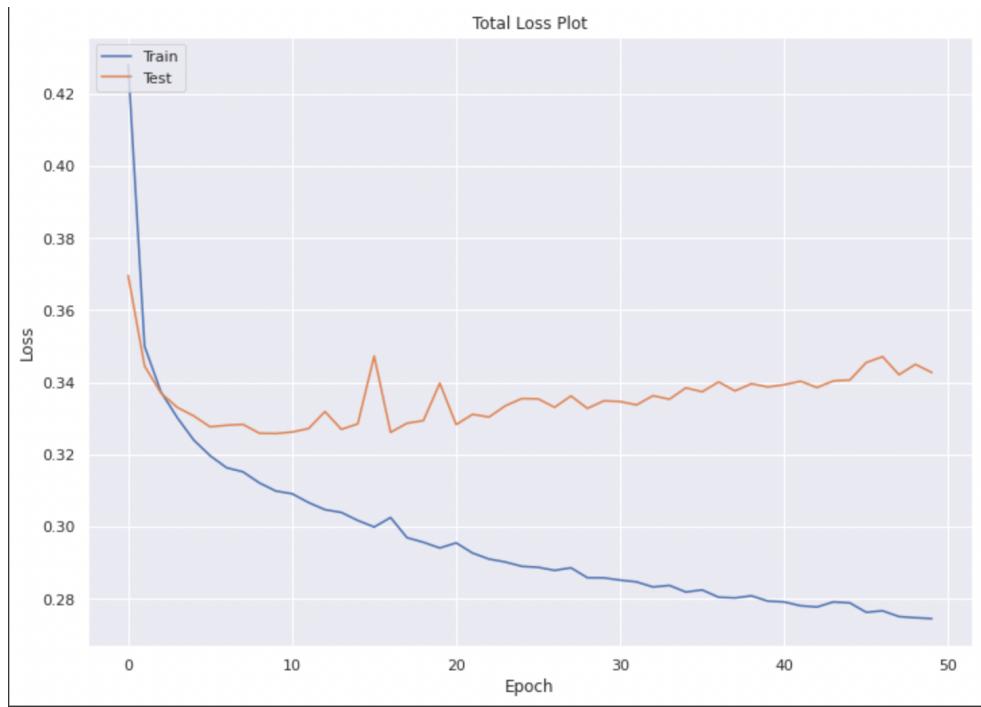
Model: "sequential_11"		
Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 48)	4656
dense_34 (Dense)	(None, 24)	1176
dense_35 (Dense)	(None, 2)	50
<hr/>		
Total params: 5,882		
Trainable params: 5,882		
Non-trainable params: 0		
<hr/>		

5. Provide graphs that compares test and training accuracy on the same plot, test and training loss on the same plot. Thus, in total two graphs with a clear labeling.

- *Model Accuracy Plot*



- *Total Loss Plot:*



## Part II – Optimizing NN

1. Include all 3 tables with different NN setups.

Dropout Tuning-

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.2	84.24%	0.4	82.20%	0.8	75.55%
Optimizer	Adam		Adam		Adam	
Activation Function	relu		relu		relu	
Initializer	-		-		-	

Optimizer Tuning-

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.2		0.2		0.2	
Optimizer	RMSprop	84.47%	SGD	83.31%	Adagrad	82.10%
Activation Function	relu		relu		relu	
Initializer	-		-		-	

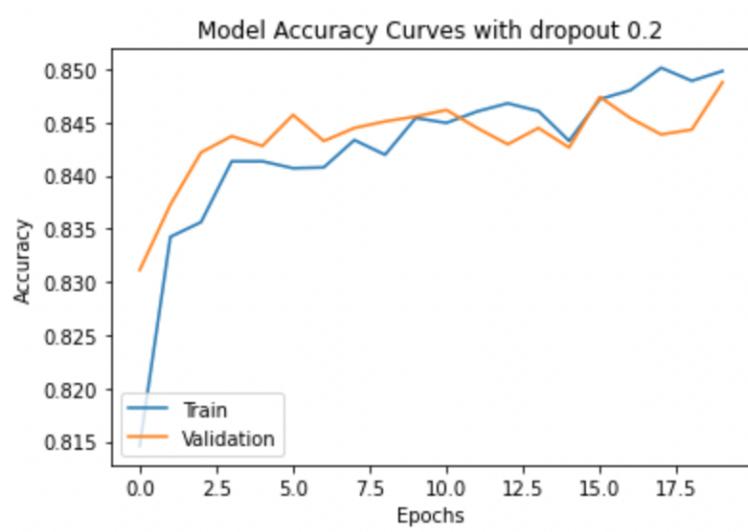
Activation Function Tuning-

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.2		0.2			
Optimizer	SGD		SGD			
Activation Function	Sigmoid	83.01	tanh	83.27	Relu	24.48
Initializer	-		-		-	

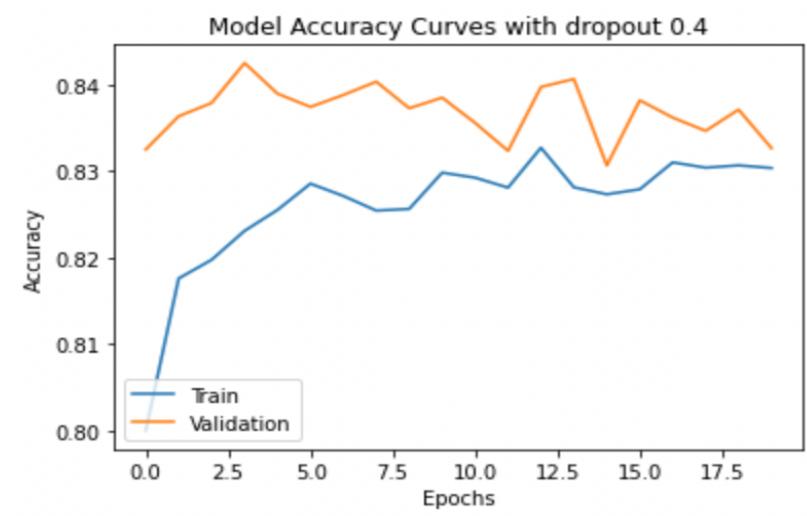
Initializer Tuning-

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.2		0.2		0.2	
Optimizer	SGD		SGD		SGD	
Activation Function	relu		relu		relu	
Initializer	Random normal	75.51%	Zeros weight	75.51%	ones	75.51

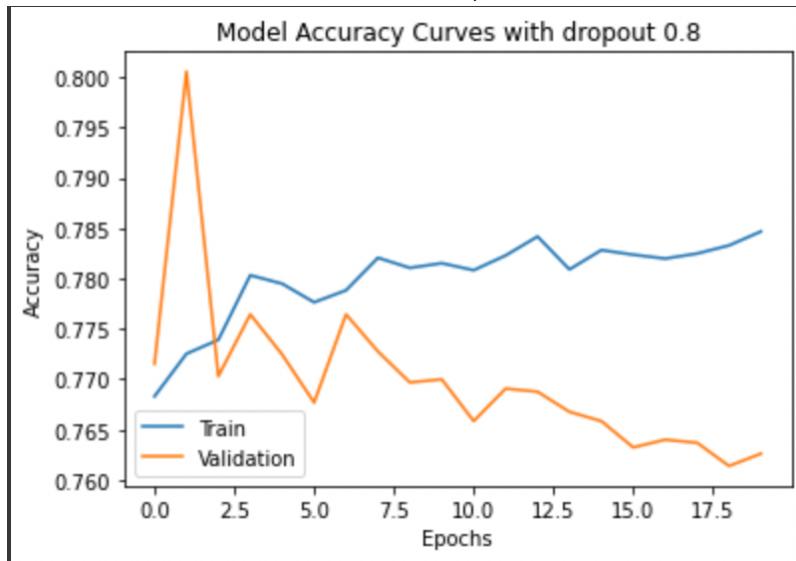
2. Provide graphs that compares test and training accuracy on the same plot for all your setups and add a short description for each graph.
  - Dropout Tuning
    1. Dropout Value=0.2 – Here for a dropout value of 0.2, the accuracy obtained is 84.24%.



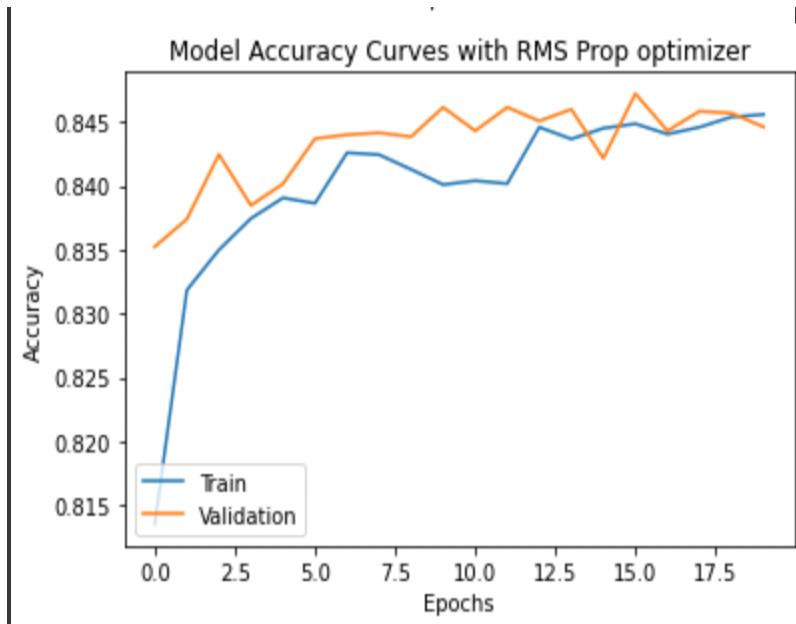
2. Dropout value=0.4 – Here for a dropout value of 0.4, the accuracy is still significant at 82.20%



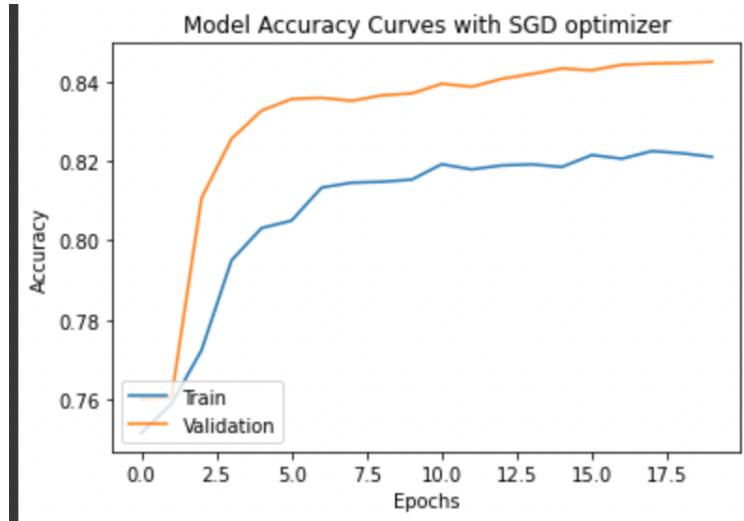
- Dropout value=0.8 - If we drop 80% of the neurons the complexity of neural network decreases a lot and as a result performance takes a massive hit



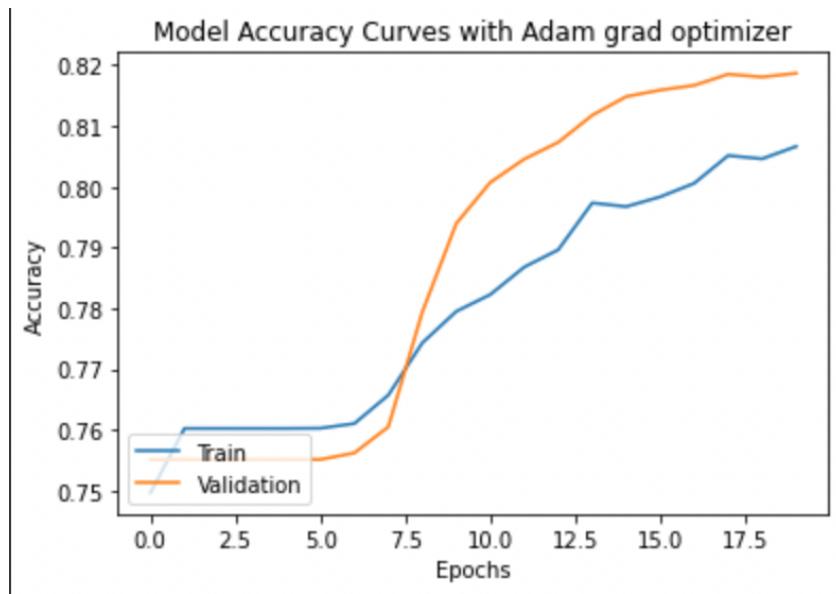
- Optimizer Tuning
  - RMSprop – With the RMSprop optimizer we achieve a maximum accuracy of 85.62%.



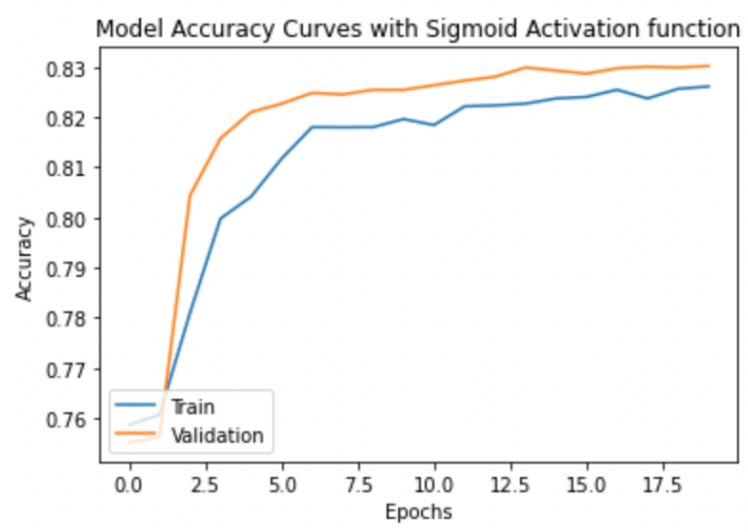
2. SGD - With Stochastic Gradient Descent optimizer the loss and accuracy graph are less jagged and asymptotic in terms of convergence



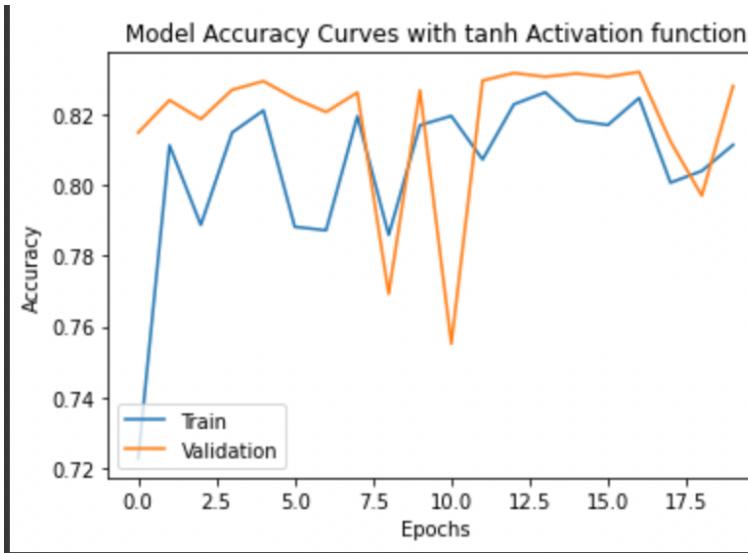
3. Adagrad - Here with Adagrad we have exponential decay of loss or convergence



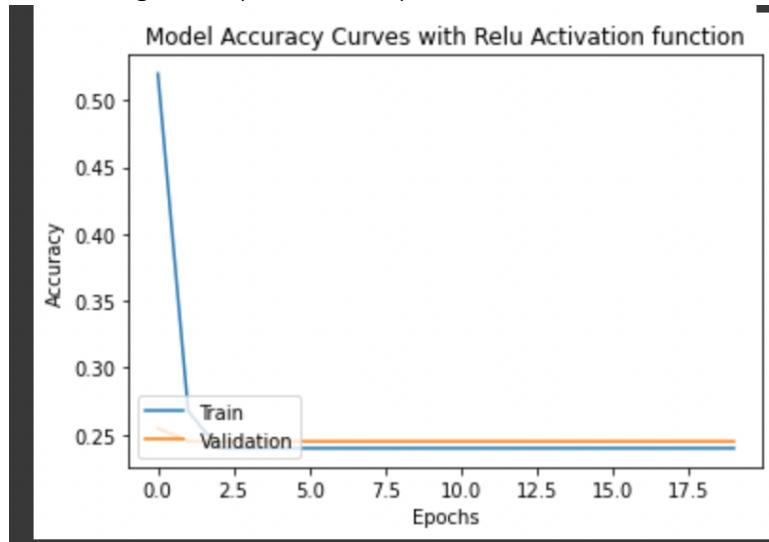
- Activation Function Tuning
  1. Sigmoid – With the sigmoid activation function we have an accuracy of around ~84.44%



2. Tanh – With tanh activation function we get a jagged plot with a maximum accuracy of 83.4%

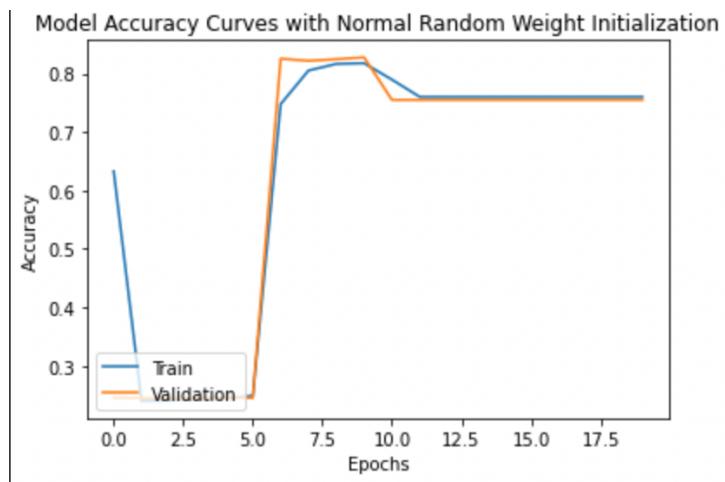


3. Relu – We get a very low accuracy of 23.5% with Relu Activation function

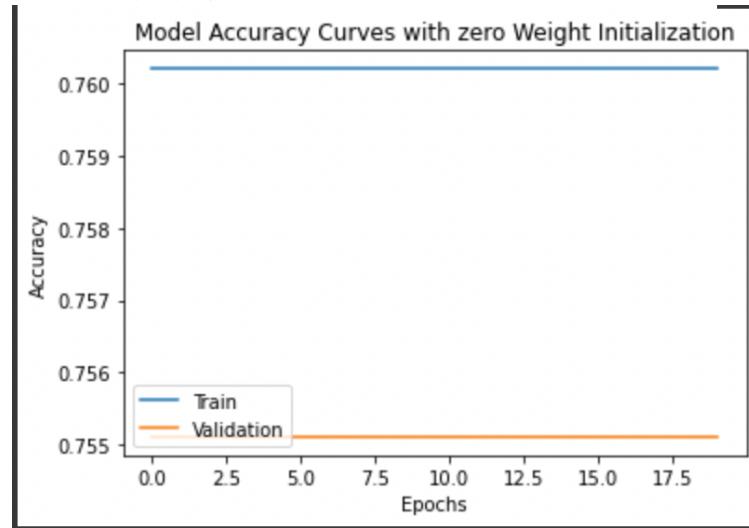


- Initializer Tuning

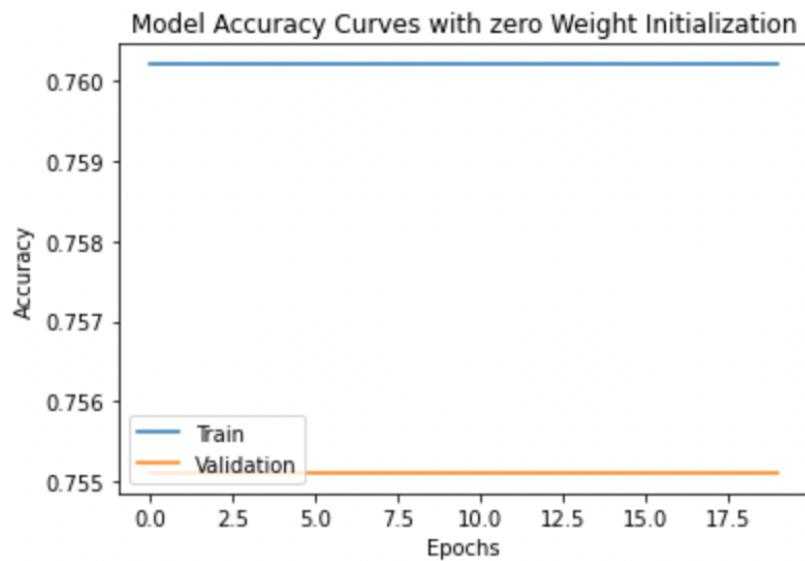
1. Random normal – We are getting an accuracy of 76% with random weight initialization



2. Zeros weight - With zero initialization there is no learning as there is no gradient to be backpropagated



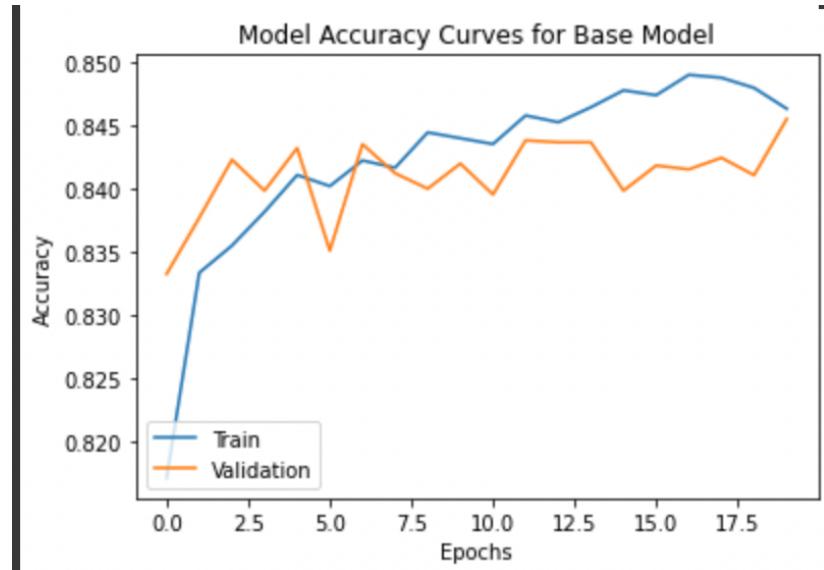
3. Ones



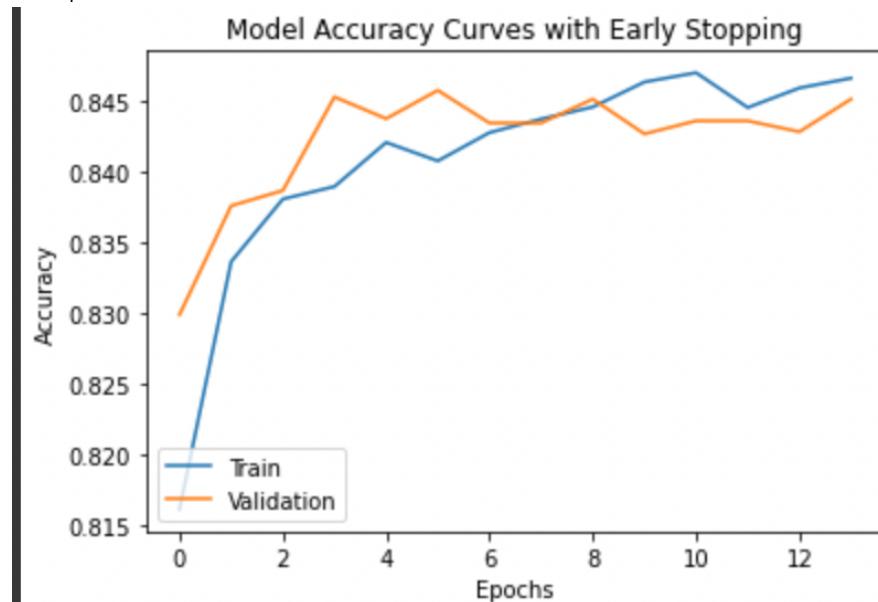
3. Provide a detailed analysis and reasoning about the NN setups that you tried.
- Hidden Layers: As the requirement was to generate a model with hidden layers no more than 3, we decided to have 2 hidden layers. The reason being with 3 hidden layers there was a performance drop without any improvement in the accuracy as compared with a model with 2 hidden layers.
  - Dropout Tuning: The dropout tuning offers an effective regularization method to reduce overfitting and improve generalization error in the deep neural networks. Here in our model, we are using dropouts to have a few different network architectures by randomly dropping out nodes during the training. It is common practice to have a dropout value in the range 0-0.5 to achieve an effective model. However, as the dropout value crosses 0.5 most of the neurons would be dropped from the network and as a result the accuracy would low.
    1. The dropout values we tried are 0.2, 0.4 and 0.8. The accuracies obtained are as 84.44%, 82.22% and 72% respectively.
    2. Here we can observe a huge drop in the accuracy with dropout value of 0.8
  - Optimizer Tuning: Optimizers are the classes that are used to change the attributes of our model such as weights and learning rate to reduce the losses. Optimizers help to get results faster.
    1. In our setups, we are using three different optimizers namely RMSProp, Stochastic Gradient Descent and Adagrad.
    2. For each of the optimizers the accuracy obtained are: 84%, 82% and 83% respectively. Hence, the RMSProp provides us the highest accuracy values.
  - Activation Function Tuning: The activation function is a link between the input of the neuron and its output to the next layer. Our activation functions have direct impact on the accuracies and the loss obtained.
    1. Here we are utilizing the following activation functions: Sigmoid, tanh and relu.
    2. We are obtaining a maximum accuracy for the model using the sigmoid function which ranges between 83-85%.
  - Initializer Tuning: Initializers are used to define initial weight values in the keras models.
    1. In our setups we have used the following initializers: Random Weights, Zero Weights and One weights.
    2. We are achieving minimum accuracy in case of zero and one initialers as there is no gradient to be back propagated.

4. Briefly discuss all the methods you used that help to improve the accuracy or training time. Provide accuracy graphs and your short descriptions.

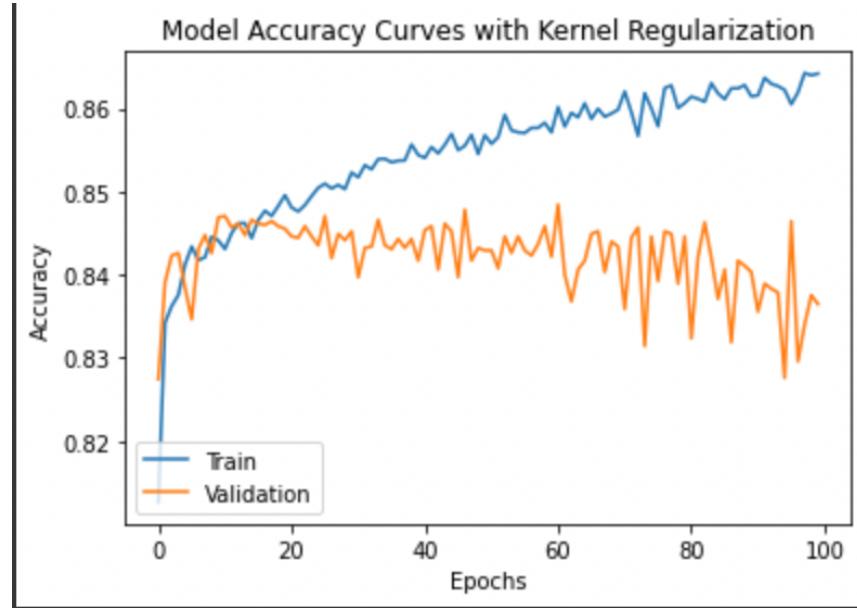
- Base Model – We are choosing the model with dropout value = 0.2 as our base model



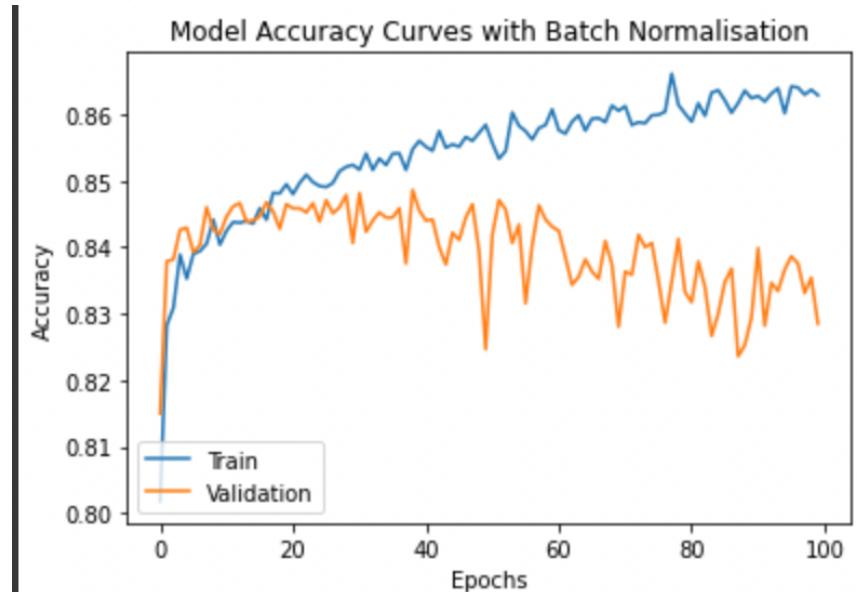
- Early Stopping: Early stopping is a regularization technique that helps with overfitting. The Early stopping feature stops the training when a specific condition has been met. In our case, we are stopping the model training when the '**val\_loss**' parameter is maximum for 5 continuous epochs.



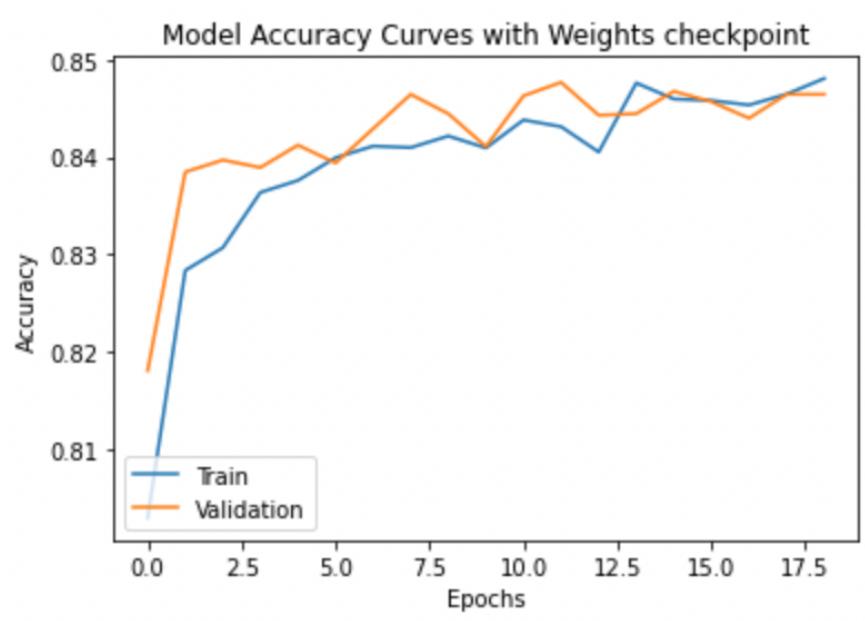
- Regularization L1 and L2- The regularization applies penalties on the layer's kernel during optimization. The penalties are then summed into the loss function that is optimized. Here we are using both L1 and L2 class regularizes.



1. Batch Normalization – The batch normalization is a layer that normalizes the input data. The layer essentially applies a transformation that maintains the mean output near to 0 and standard deviation near to 1.



- Weights Checkpoint – The checkpoint function is used to store the weights values in a temporary location when a specific condition is met. In our case we are storing the weights when the val\_accuracy is maximum. Once the weights have been saved, they would be used during the evaluation to achieve a maximum accuracy.



## Part III: Building a CNN

1. Provide brief details about the nature of your dataset. What is it about? What type of data are we encountering? How many entries and variables does the dataset comprise? Provide the main statistics about the entries of the dataset.
  - Fashion-MNIST is a dataset of clothing images. The dataset consists of a training set of 60,000 images and a test set of 10,000 images. Each of the images has a dimension of 28x28 and is linked with one of the 10 classes. The 10 classes have a label associated as below:

Class Label Number	Class Label Value
0	T-shirt
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle Boot

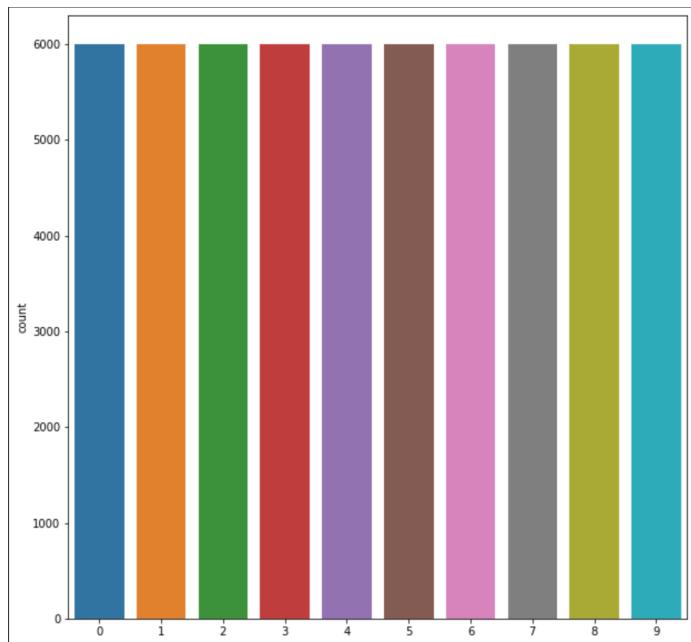
- The fashion-mnist is a suitable and challenging replacement for MNIST dataset and it can be used to build models with different neural architecture setup.

2. Provide at least 3 visualization graphs with short description for each graph.

- The below visualization displays the grey-scale images of the fashion-mnist dataset before preprocessing.



- The below visualization displays the count of each of the class in the fashion-mnsit dataset which indicates a cumulative value of 60000 class of images in the training dataset.



- The below image displays the fashion-mnist images after preprocessing and normalization.



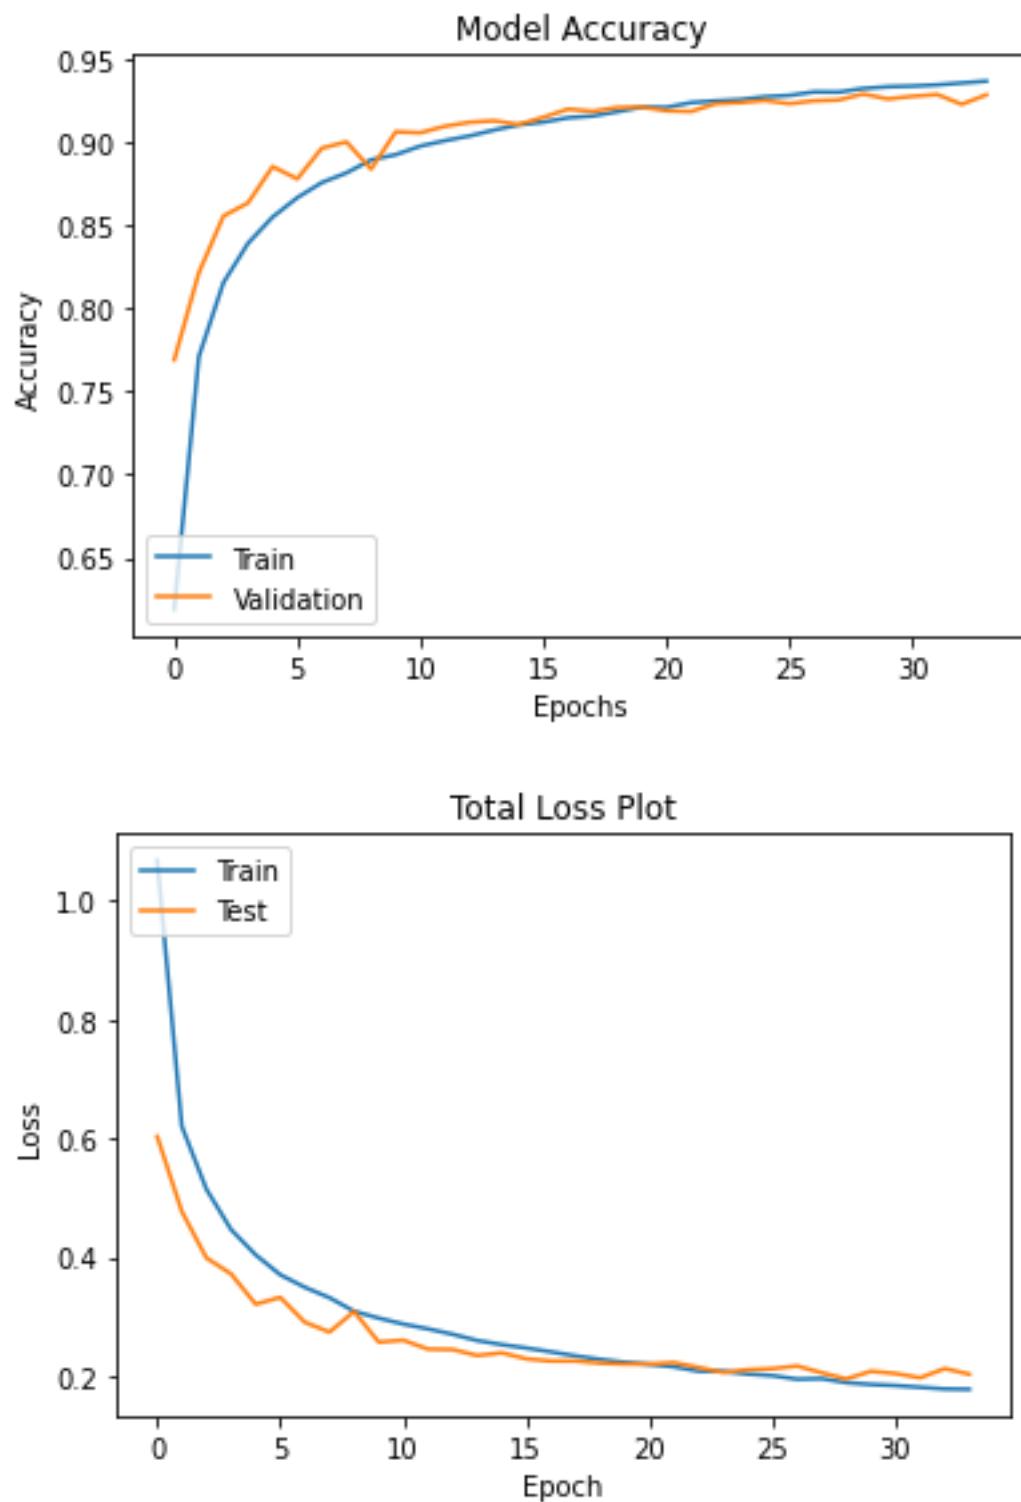
3. Provide the architecture structure of your CNN.

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (BatchN ormalization)	(None, 28, 28, 32)	128
conv2d_1 (Conv2D)	(None, 26, 26, 32)	9248
batch_normalization_1 (BatchN ormalization)	(None, 26, 26, 32)	128
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 13, 13, 64)	18496
batch_normalization_2 (BatchN ormalization)	(None, 13, 13, 64)	256
conv2d_3 (Conv2D)	(None, 11, 11, 64)	36928
batch_normalization_3 (BatchN ormalization)	(None, 11, 11, 64)	256
max_pooling2d_1 (MaxPooling 2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
conv2d_4 (Conv2D)	(None, 5, 5, 128)	73856
batch_normalization_4 (BatchN ormalization)	(None, 5, 5, 128)	512
conv2d_5 (Conv2D)	(None, 3, 3, 128)	147584
batch_normalization_5 (BatchN ormalization)	(None, 3, 3, 128)	512
max_pooling2d_2 (MaxPooling 2D)	(None, 1, 1, 128)	0
dropout_2 (Dropout)	(None, 1, 1, 128)	0
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 256)	33024
dense_1 (Dense)	(None, 10)	2570
<hr/>		
Total params:	323,818	
Trainable params:	322,922	
Non-trainable params:	896	

4. Discuss how the improvement tools works on CNN architectures.

- *Earlystopping*: The earlystopping option which we utilized here stops the model training when the ‘val\_loss’ is minimum for 5 consecutive epochs. The model won’t stop training if the loss value is reducing further, however as soon as the loss stabilizes i.e., there is no decrease in the loss value the model stops training. The earlystopping is used in combination with the weights checkpoint tool.
- *Weights Checkpoint*: The purpose of the weights checkpoint tool is that it saves the weight values in a temporary location whenever a maximum val\_accuracy value is obtained during the training. As soon as the training stops due to the earlystopping tool, the saved weights are used for the evaluation.

5. Provide graphs that compares test and training accuracy on the same plot, test and training loss on the same plot. Thus in total two graph with a clear labeling.



## Part IV: Optimizing CNN + Data Augmentation

1. Include all 3 tables with different CNN setups.

- Dropout Tuning:

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.15,0.20,0.25	93.01%	0.10,0.20,0.30	92.10%	0.1,0.1,0.1	92.43%
Optimizer	adam		adam		adam	
Activation Function	relu		relu		relu	
Initializer	Glorot Normal		Glorot Normal		Glorot Normal	
Kernel size	(3, 3)		(3, 3)		(3, 3)	

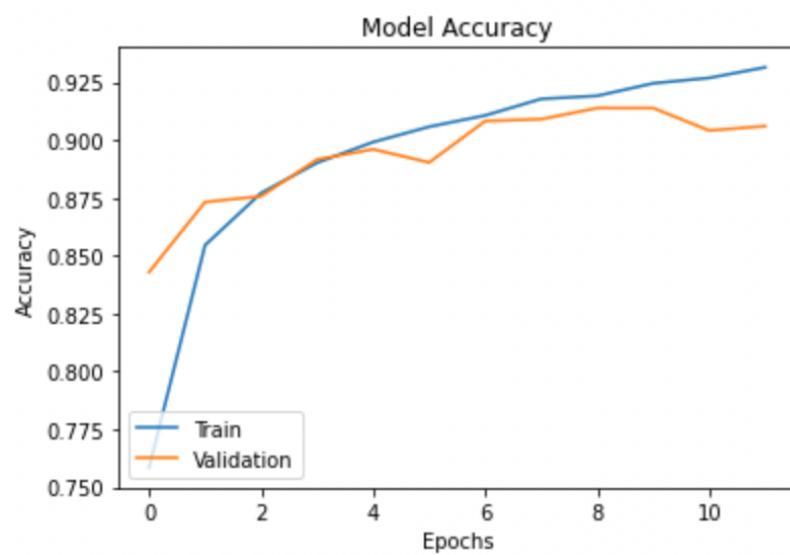
- Optimizer Tuning:

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.15,0.20,0.25		0.15,0.20,0.25		0.15,0.20,0.25	
Optimizer	adam	92.57%	RMS prop	92.03%	Adagrad	89.86%
Activation Function	relu		relu		relu	
Initializer	Glorot Normal		Glorot Normal		Glorot Normal	
Kernel size	(3, 3)		(3, 3)		(3, 3)	

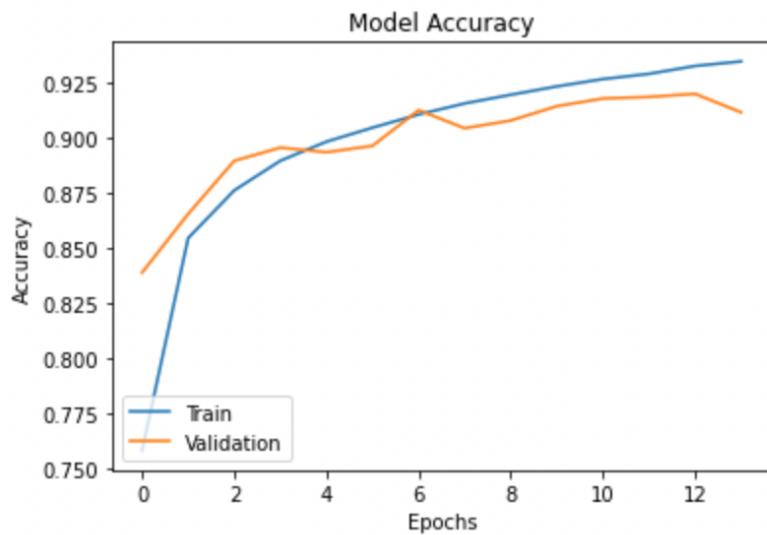
- Kernel Size Tuning:

	Setup 1	Accuracy	Setup 2	Accuracy	Setup 3	Accuracy
Dropout	0.15,0.20,0.25		0.15,0.20,0.25		0.15,0.20,0.25	
Optimizer	adam		adam		adam	
Activation Function	relu		relu		relu	
Initializer	Glorot Normal		Glorot Normal		Glorot Normal	
Kernel size	(3, 3)	92.77%	(5, 5)	92.05%	(7, 7)	90%

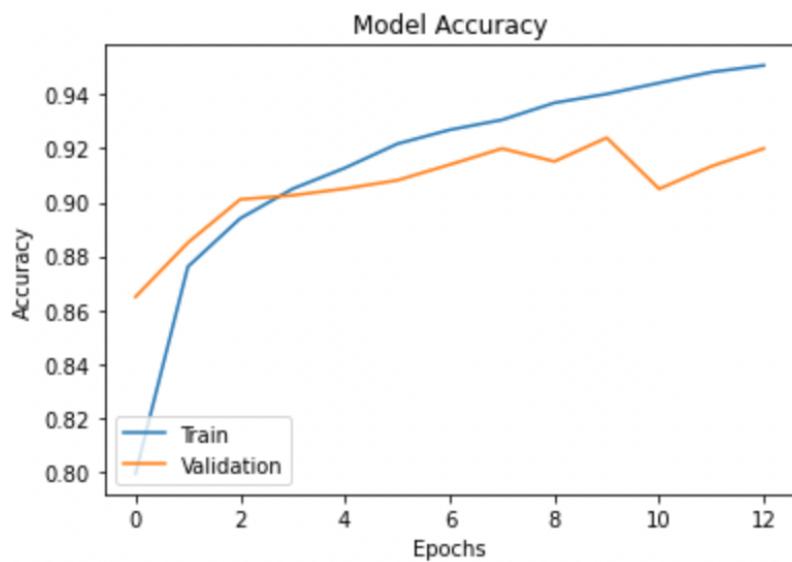
- Provide graphs that compares test and training accuracy on the same plot for all your setups and add a short description for each graph.
  - Dropout Tuning (0.15, 0.20, 0.25) – The accuracy obtained with these dropout values is 90.06%



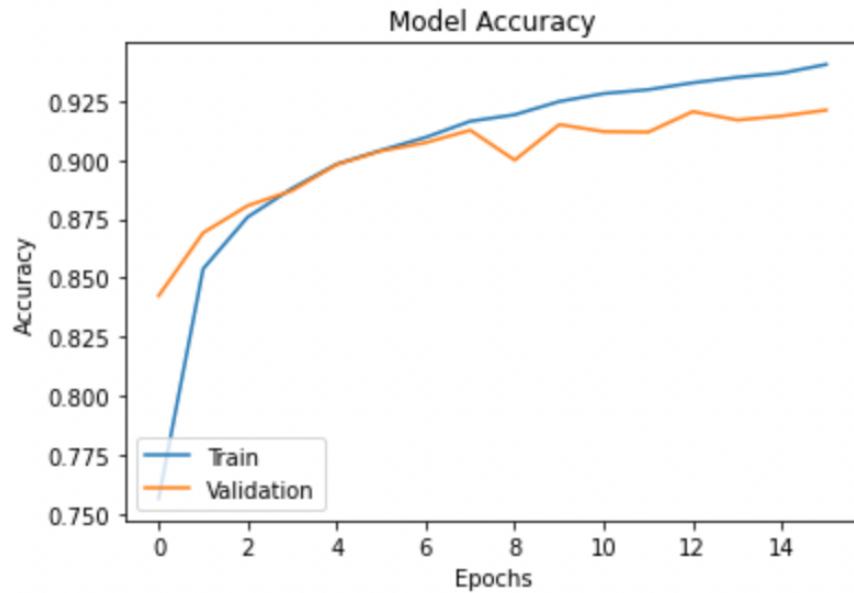
- Dropout Tuning (0.10, 0.20, 0.30) - The accuracy obtained with these dropout values is 91.17%



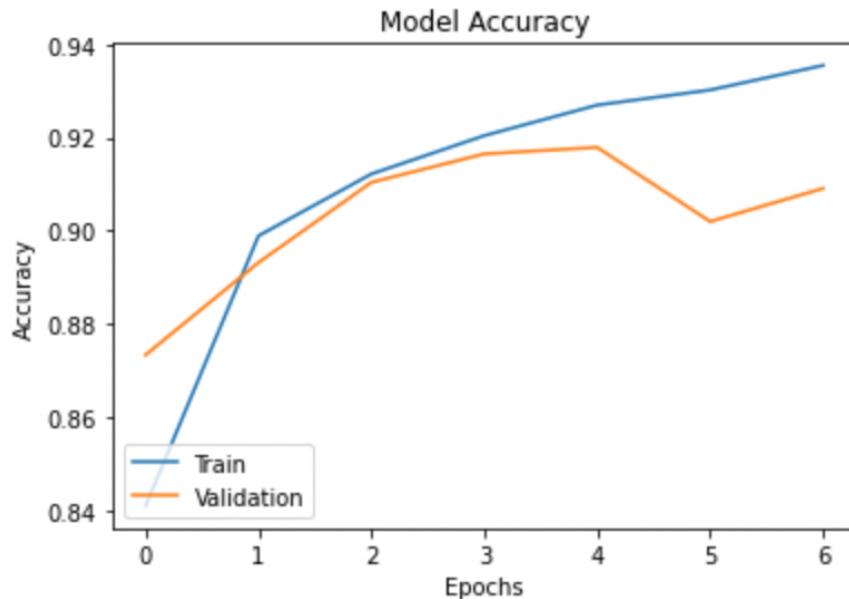
- Dropout Tuning (0.1, 0.1, 0.1) - The accuracy obtained with these dropout values is 92.0%



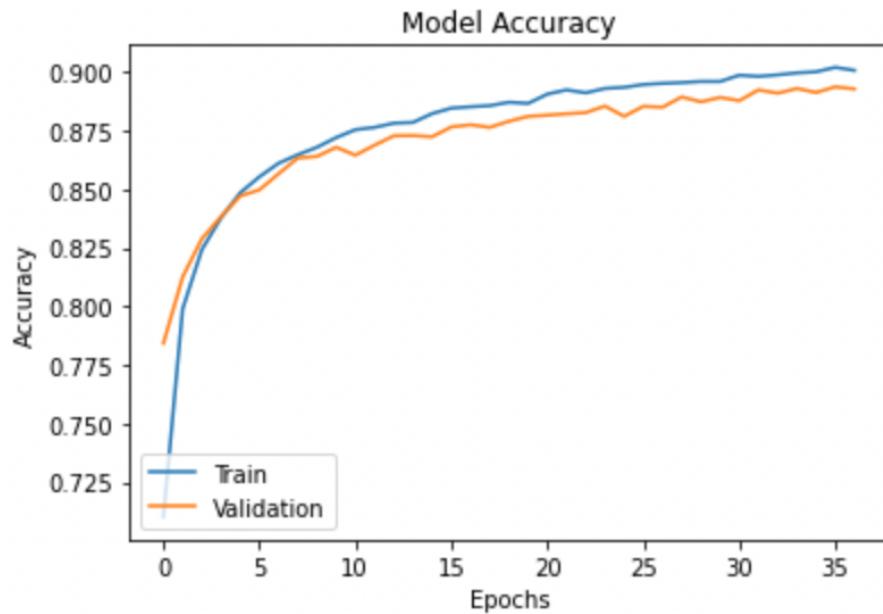
- Optimizer Tuning (Adam) - The accuracy obtained Adam optimizer is 92.1%



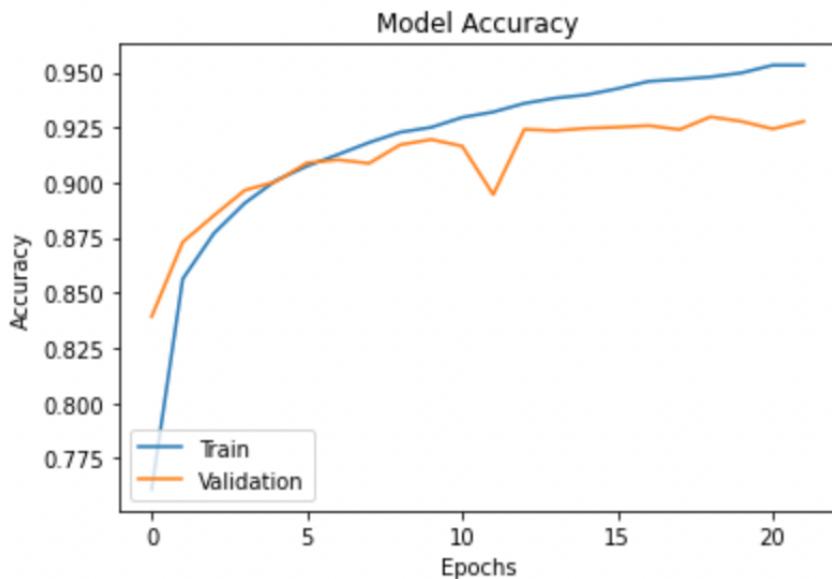
- Optimizer Tuning (RMSProp) - The accuracy obtained Adam optimizer is 90.91%



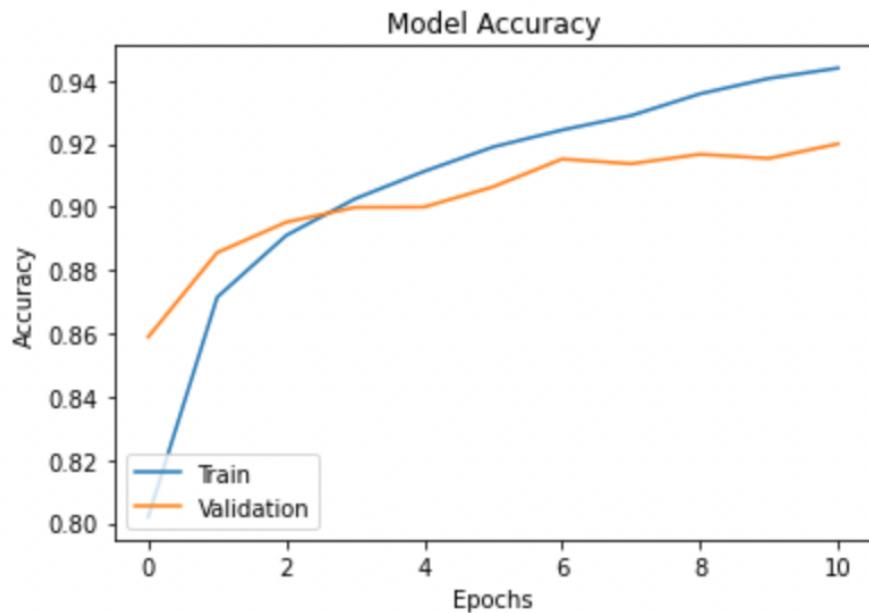
- Optimizer Tuning (Adagrad) - The accuracy obtained Adam optimizer is 89.29%



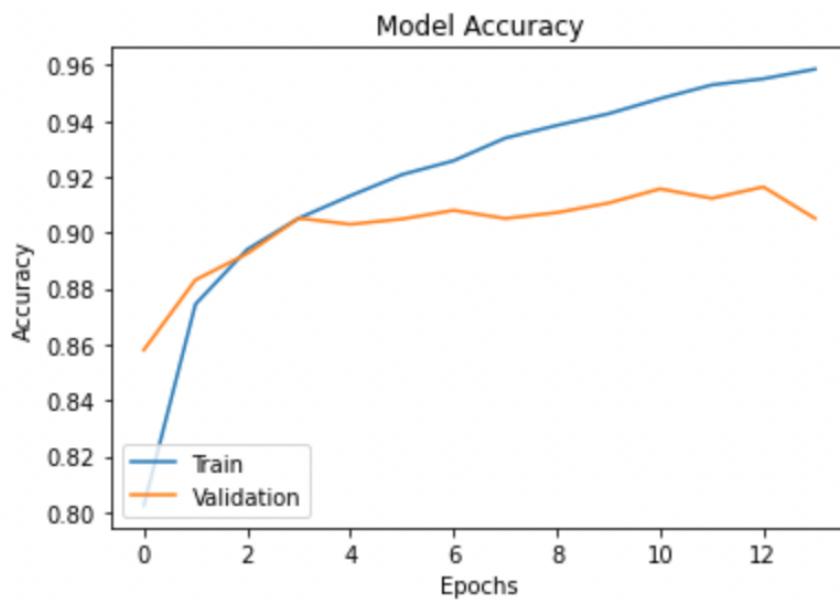
- Kernel Size Tuning (3, 3) - The accuracy obtained with kernel\_size=(3,3) is 92.7%



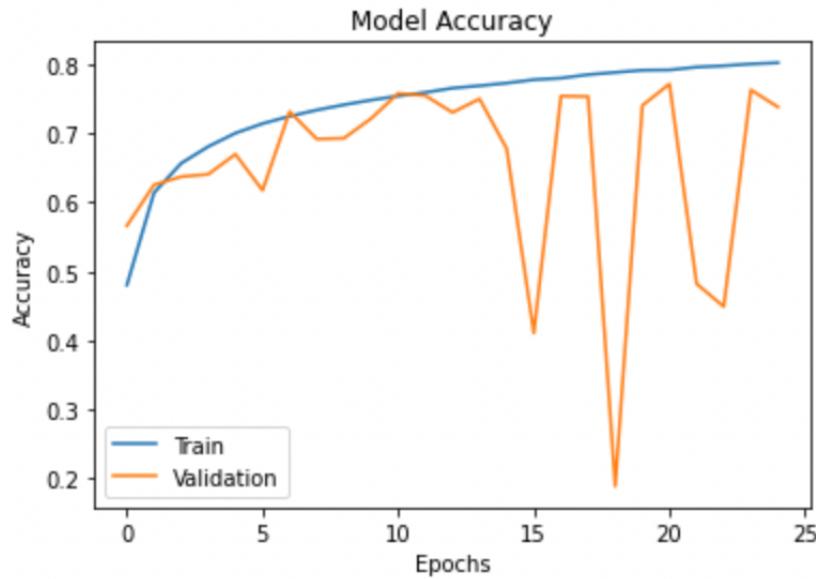
- Kernel Size Tuning (5, 5) - The accuracy obtained with `kernel_size=(5,5)` is 92.0%



- Kernel Size Tuning(7, 7) - The accuracy obtained with `kernel_size=(7,7)` is 90.5%



- Data Augmentation:



3. Provide a detailed analysis and reasoning about the CNN setups that you tried.
- Dropout Tuning: The dropout tuning offers an effective regularization method to reduce overfitting and improve generalization error in the deep neural networks. Here in our model, we are using dropouts to have a few different network architectures by randomly dropping out nodes during the training. It is common practice to have a dropout value in the range 0-0.5 to achieve an effective model. However, as the dropout value crosses 0.5 most of the neurons would be dropped from the network and as a result the accuracy would low.
    1. In this setup we tried different set of dropout values for each layer.
    2. The values we tried are: (0.15, 0.20, 0.25), (0.10, 0.20, 0.30) and (0.1, 0.1, 0.1).
    3. We obtained a maximum accuracy of 92% when the dropout value was same and minimum across all the layers. When we tried a dropout value of 0.2 across all the layers even though we were achieving a maximum accuracy of 93%, it was not consistent during the runs.
  - Optimizer Tuning: Optimizers are the classes that are used to change the attributes of our model such as weights and learning rate to reduce the losses. Optimizers help to get results faster.

1. In our setups, we are using three different optimizers namely RMSProp, Adam and Adagrad.
  2. For each of the optimizers the accuracy obtained are: 90.91%, 92.1% and 89.92% respectively. Hence, the adagrad optimizer provides us the lowest accuracy values whereas the Adam optimizer provides the highest.
- Kernel Size Tuning: The kernel size in a keras model indicates the height and width of the convolution model.
    1. In our setups, we are using three different kernel size values: (3,3), (5,5) and (7,7).
    2. Even though there is no direct relationship between kernel size and the accuracy. The kernel size setup which are trying here is to check if using a kernel of specific size covers majority of the features. A larger kernel size may detect the features missed by the smaller kernel size and vice-versa.
    3. Here we are achieving the accuracy of the values: 92.77%, 92.0% and 90.05%
4. Briefly discuss methods that you used for data augmentation and reason how it helped to increase the accuracy of the model.
    - We used the following data augmentation methods:

Horizontal Flip
Vertical Flip
90 Degree Rotation
Zoom

- The data augmentation process is generally utilized when the input dataset size is insufficient for training the model, so using the augmentation process increases the size of the data.

## BONUS QUESTION

1. The accuracy obtained using the perceptron built from scratch is 52.23%
2. The accuracy obtained using the Logistic Regression is 88.05%

*CONTRIBUTION:*

<b>Team Member</b>	<b>Assignment Part</b>	<b>Contribution</b>
Sachin Sarsambi	Part 1	50%
Naman Tejaswi		50%
Sachin Sarsambi	Part 2	50%
Naman Tejaswi		50%
Sachin Sarsambi	Part 3	50%
Naman Tejaswi		50%
Sachin Sarsambi	Part 4	50%
Naman Tejaswi		50%

*References:*

- 1) <https://keras.io/api/>
- 2) <https://keras.io/api/models/sequential/>
- 3) <https://numpy.org/>
- 4) [https://keras.io/api/models/model\\_training\\_apis/#fit-method](https://keras.io/api/models/model_training_apis/#fit-method)
- 5) [https://keras.io/api/callbacks/model\\_checkpoint/](https://keras.io/api/callbacks/model_checkpoint/)
- 6) <https://keras.io/api/optimizers/>
- 7) <https://keras.io/api/layers/activations/>
- 8) [https://keras.io/api/callbacks/early\\_stopping/](https://keras.io/api/callbacks/early_stopping/)
- 9) [https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)